

SARSA Algorithm

State-Action-Reward-State-Action (SARSA) is model-free Reinforcement Learning algorithm used to find the optimum policy for a Markov Decision Process (MDP). It was first proposed in "Online Q-Learning using Connectionist Systems" by (Rummery and Niranjan, 1994) It is a member of the Temporal-Difference (TD) class of algorithms similar to Q-Learning and is an on-policy algorithm, meaning that it requires an initial policy to iterate on.

SARSA often uses ϵ -greedy (epsilon-greedy) as its initial policy as it provides a balance between exploration and exploitation. According to (Sutton and Barto, 1998), SARSA converges with probability 1 to an optimal policy π^* and action-value function $Q(a, s)$ after all state-action pairs are visited an infinite number of times.

Algorithm

Firstly, for all state and action pairs (s, a) , initialise the action-value function $Q(s, a)$ with any number. From the initial state, the next action is chosen by the policy π . For a given state s , $Q(s, a)$ is be updated using equation 1. Figure 1 shows the states, rewards and actions used in the equation.

Where:

- r is the immediate reward after taking action a in state s .
- γ is the discount factor where $[0 \leq \gamma \leq 1]$
- s_{t+1} is the next state.
- a_{t+1} is the next action according to the policy.
- α is the learning rate $0 \leq \alpha \leq 1$

$$Q(s_t, a_t) = Q(s_t, a_t) + \alpha[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)] \quad (1)$$




Figure 1 – State transitions (Sutton and Barto, 1998 pg. 145)

The current state s is then shifted to state s_{t+1} . This process is repeated for every state transition until a termination state is reached or all Q-values converge. As the algorithm continues, $Q(s, a)$ converges to $Q^*(s, a)$ and the optimal policy π^* is found.

The pseudo-code for SARSA is shown in Figure 2 provided by (Sutton and Barto, 1998, pg 143).

```
Initialize  $Q(s, a)$  arbitrarily
Repeat (for each episode):
  Initialize  $s$ 
  Choose  $a$  from  $s$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
  Repeat (for each step of episode):
    Take action  $a$ , observe  $r, s'$ 
    Choose  $a'$  from  $s'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
     $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$ 
     $s \leftarrow s'; a \leftarrow a'$ 
  until  $s$  is terminal
```

Figure 2 – SARSA pseudo-code

References

Rummery, G. A. and Niranjan, M. (1994) *Online Q-Learning using Connectionist Systems*. : Cambridge University Engineering Department

Sutton, R.S. and Barto, A.G. (1998) *Reinforcement Learning: An Introduction*. 2nd ed. : Mit Press.